

Stat 6390.001: High Dimensional Probability and Statistics, Fall 2026

Topic 1

High Dimensional Probability and Statistics

Lecturer: Yun Wei

Scribe: Students in Fall 2026

Disclaimer: *These notes have not been subjected to the usual scrutiny reserved for formal publications. They may be distributed outside this class only with the permission of the Instructor.*

Contents

1	Basic tail and concentration bounds	5
1.1	Introduction	5
1.1.1	One-dimensional and multivariate CLT	5
1.2	Classical Tail Bounds	6
1.2.1	Markov and Chebyshev inequalities	6
1.2.2	From Markov to the Chernoff bound	7
1.3	Sub-Gaussian Random Variables	7
1.3.1	Gaussian Tail Bound	7
1.3.2	Definition and basic tail bounds	8
1.3.3	Example of Sub-Gaussian Variables	10
1.3.4	Hoeffding’s Inequality	10
1.4	Symmetrization	13
1.5	Sub-exponential Random Variables	14
1.5.1	Bernstein’s Condition	15
1.5.2	Sums of Independent Sub-exponential Random Variables	17
1.5.3	Application: The Johnson–Lindenstrauss Lemma	18
1.6	Martingale-Based Methods	20
1.6.1	Azuma–Hoeffding Inequality	21
1.6.2	Bounded Differences	22
1.7	Lipschitz Functions of Gaussian Variables	25
2	Concentration of measure	27
2.1	Concentration by Entropic Techniques	27
2.2	new	28
2.2.1	Optimal Transport Theory	28
2.2.2	Kullback–Leibler Divergence	30
3	Uniform Laws of Large Numbers	33
3.1	Motivation	33
3.1.1	Uniform Convergence of CDFs	33
3.1.2	Uniform Laws for General Function Classes	34

Chapter 1

Basic tail and concentration bounds

1.1 Introduction

A central challenge in high-dimensional statistics is the *curse of dimensionality*. As the dimension d grows, the sample space expands rapidly, so a fixed number of observations becomes increasingly sparse relative to the size of the space. As a result, methods and intuition that work well in low dimensions may become much less informative in high-dimensional settings.

High-dimensional problems also exhibit geometric and probabilistic behavior that can differ substantially from low-dimensional intuition. For example, if

$$X \sim N(0, I_d),$$

then the density is maximized at the origin, but a typical sample is not concentrated near the origin. Instead,

$$\|X\|_2 \approx \sqrt{d},$$

so most samples lie near a thin shell of radius approximately $\sqrt{d}$.

These phenomena motivate the study of probabilistic tools that remain useful when the dimension is large. In particular, this lecture develops tail and concentration bounds that quantify how likely a random variable is to deviate from its mean, with an emphasis on non-asymptotic guarantees.

1.1.1 One-dimensional and multivariate CLT

Let $X_1, \dots, X_n$ be i.i.d. random variables with mean μ and finite variance σ^2 . The one-dimensional CLT states that

$$\sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n X_i - \mu \right) \xrightarrow{d} N(0, \sigma^2) \quad \text{as } n \rightarrow \infty. \quad (1.1)$$

The multivariate version has the same form. Suppose now that $X_1, \dots, X_n$ are i.i.d. random vectors in $\mathbb{R}^d$ with

$$\mathbb{E}[X_i] = \mu \in \mathbb{R}^d, \quad \text{Cov}(X_i) = \Sigma.$$

Then

$$\sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n X_i - \mu \right) \xrightarrow{d} N(0, \Sigma) \quad \text{as } n \rightarrow \infty. \quad (1.2)$$

For a d -dimensional Gaussian vector $X \sim N(\mu, \Sigma)$ with positive definite covariance matrix Σ , the density is

$$f(x) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}. \quad (1.3)$$




Figure 1.1: A schematic illustration of a high-dimensional standard Gaussian distribution.

Example 1.1.1 (A First High-Dimensional Phenomenon). Consider the standard Gaussian case $X \sim N(0, I_d)$. The density is maximized at the origin, so the mode is $x = 0$. However, when d is large, a typical draw is not close to the origin. Indeed,

$$\mathbb{E}\|X\|_2^2 = \mathbb{E}\left[\sum_{j=1}^d X_j^2\right] = \sum_{j=1}^d \mathbb{E}[X_j^2] = d. \quad (1.4)$$

Thus the typical Euclidean radius is of order $\sqrt{d}$; more precisely, $\|X\|_2^2$ is a chi-square random variable with d degrees of freedom and is concentrated around d . Hence most of the Gaussian probability mass lies in a thin shell with radius approximately $\sqrt{d}$, even though the density itself is largest at the origin, shown in Figure 1. This illustrates a basic feature of high-dimensional geometry: the location of the mode need not describe where a typical sample lies.

1.2 Classical Tail Bounds

The purpose of a tail bound is to control probabilities such as $\mathbb{P}(X \geq t)$ or $\mathbb{P}(|X - \mathbb{E}X| \geq t)$. The amount of information available about the distribution determines the quality of the bound. Markov's inequality uses only a first moment; Chebyshev's inequality uses a second moment; and the Chernoff method exploits the moment generating function (MGF), thereby incorporating information from all moments when the MGF exists [1].

1.2.1 Markov and Chebyshev inequalities

Proposition 1.2.1 (Markov's inequality). *If random variable X is nonnegative and $\mathbb{E}[X] < \infty$, then for*

every $t > 0$,

$$\mathbb{P}(X \geq t) \leq \frac{\mathbb{E}[X]}{t}. \quad (1.5)$$

Proposition 1.2.2 (Chebyshev's inequality). *If random variable X has finite variance and $\mu = \mathbb{E}[X]$, then for every $t > 0$,*

$$\mathbb{P}(|X - \mu| \geq t) \leq \frac{\text{Var}(X)}{t^2}. \quad (1.6)$$

Chebyshev's inequality is an immediate consequence of Markov's inequality applied to the nonnegative random variable $(X - \mu)^2$. More generally, if the k th centered absolute moment exists, then

$$\mathbb{P}(|X - \mu| \geq t) = \mathbb{P}(|X - \mu|^k \geq t^k) \leq \frac{\mathbb{E}|X - \mu|^k}{t^k}. \quad (1.7)$$

1.2.2 From Markov to the Chernoff bound

The exponential function is particularly effective because it converts sums into products and leads naturally to the moment generating function.

Lemma 1.2.3 (Chernoff bound). *Let $\mu = \mathbb{E}[X]$, and suppose*

$$\phi(\lambda) = \mathbb{E}[e^{\lambda(X - \mu)}]$$

exists for $\lambda \in [0, b]$. Then for every $t > 0$,

$$\log \mathbb{P}(X - \mu \geq t) \leq \inf_{\lambda \in [0, b]} \{\log \phi(\lambda) - \lambda t\}. \quad (1.8)$$

Proof. For $\lambda > 0$. Since the exponential function is increasing,

$$\mathbb{P}(X - \mu \geq t) = \mathbb{P}(\lambda(X - \mu) \geq \lambda t) \quad (1.9)$$

$$= \mathbb{P}(e^{\lambda(X - \mu)} \geq e^{\lambda t}). \quad (1.10)$$

Applying Markov's inequality to the nonnegative random variable $e^{\lambda(X - \mu)}$ gives

$$\mathbb{P}(X - \mu \geq t) \leq \frac{\mathbb{E}[e^{\lambda(X - \mu)}]}{e^{\lambda t}} = \phi(\lambda)e^{-\lambda t}. \quad (1.11)$$

Taking logarithms yields

$$\log \mathbb{P}(X - \mu \geq t) \leq \log \phi(\lambda) - \lambda t.$$

Because the inequality holds for every admissible λ , minimizing the right-hand side over $\lambda \in [0, b]$ gives the result. $\square$

Remark 1.2.4. The optimization over λ is essential. Different values of λ produce different valid bounds, and the infimum selects the sharpest one obtained from this argument. In applications, the main work is therefore to control the MGF and then solve a one-dimensional optimization problem.

1.3 Sub-Gaussian Random Variables

1.3.1 Gaussian Tail Bound

As a first example, consider a Gaussian random variable

$$X \sim N(\mu, \sigma^2).$$

Its moment generating function is

$$\mathbb{E}[e^{\lambda X}] = \exp\left(\mu\lambda + \frac{\sigma^2\lambda^2}{2}\right), \quad \lambda \in \mathbb{R}. \quad (1.12)$$

Applying the preceding lemma, the upper-tail bound is determined by

$$\inf_{\lambda>0} \left\{ \log \mathbb{E}[e^{\lambda(X-\mu)}] - \lambda t \right\} = \inf_{\lambda>0} \left\{ \frac{\sigma^2\lambda^2}{2} - \lambda t \right\}. \quad (1.13)$$

Here we used the fact that

$$X - \mu \sim N(0, \sigma^2).$$

The quadratic expression is minimized at

$$\lambda^* = \frac{t}{\sigma^2},$$

which gives

$$\inf_{\lambda>0} \left\{ \frac{\sigma^2\lambda^2}{2} - \lambda t \right\} = -\frac{t^2}{2\sigma^2}. \quad (1.14)$$

Therefore, for every $t > 0$,

$$\mathbb{P}(X - \mu \geq t) \leq \exp\left(-\frac{t^2}{2\sigma^2}\right). \quad (1.15)$$

(Normal distribution is light tail)

1.3.2 Definition and basic tail bounds

Definition 1.3.1 (Sub-Gaussian random variable). A random variable X with mean $\mu = \mathbb{E}[X]$ is *sub-Gaussian with parameter $\sigma > 0$* if

$$\mathbb{E}\left[e^{\lambda(X-\mu)}\right] \leq \exp\left(\frac{\sigma^2\lambda^2}{2}\right) \quad \text{for all } \lambda \in \mathbb{R}. \quad (1.16)$$

Remark 1.3.2. If X is σ -sub-Gaussian for some $\sigma > 0$, then X is also σ' -sub-Gaussian for some $\sigma' \geq \sigma$.

The definition says that the centered MGF of X grows no faster than the centered MGF of a Gaussian $N(0, \sigma^2)$. Therefore, the same Chernoff calculation gives the Gaussian-type upper-tail bound

$$\mathbb{P}(X - \mu \geq t) \leq \exp\left(-\frac{t^2}{2\sigma^2}\right), \quad t > 0. \quad (1.17)$$

Lemma 1.3.3 (Symmetry of sub-Gaussianity). *A random variable X is sub-Gaussian with parameter σ if and only if $-X$ is sub-Gaussian with the same parameter.*

Proof. Let $\mu = \mathbb{E}[X]$. Recall that X is sub-Gaussian with parameter σ if

$$\mathbb{E}\left[e^{\lambda(X-\mu)}\right] \leq \exp\left(\frac{\sigma^2\lambda^2}{2}\right), \quad \lambda \in \mathbb{R}. \quad (1.18)$$

Suppose first that X is sub-Gaussian. Since $\mathbb{E}[-X] = -\mu$,

$$\mathbb{E}\left[e^{\lambda((-X)-\mathbb{E}[-X])}\right] = \mathbb{E}\left[e^{-\lambda(X-\mu)}\right] \quad (1.19)$$

$$\leq \exp\left(\frac{\sigma^2(-\lambda)^2}{2}\right) \quad (1.20)$$

$$= \exp\left(\frac{\sigma^2\lambda^2}{2}\right). \quad (1.21)$$

Hence $-X$ is sub-Gaussian with parameter σ .

Conversely, suppose that $-X$ is sub-Gaussian. Then

$$\mathbb{E} \left[e^{\lambda(X-\mu)} \right] = \mathbb{E} \left[e^{-\lambda((-X)-\mathbb{E}[-X])} \right] \quad (1.22)$$

$$\leq \exp \left(\frac{\sigma^2(-\lambda)^2}{2} \right) \quad (1.23)$$

$$= \exp \left(\frac{\sigma^2\lambda^2}{2} \right). \quad (1.24)$$

Therefore X is sub-Gaussian with parameter σ . $\square$

Lemma 1.3.4 (Lower Tail). *If X is sub-Gaussian with parameter σ , then for every $t > 0$,*

$$\mathbb{P}(X - \mathbb{E}[X] \leq -t) \leq \exp \left(-\frac{t^2}{2\sigma^2} \right). \quad (1.25)$$

Proof. By the symmetry lemma, $-X$ is also sub-Gaussian with parameter σ . Moreover,

$$\mathbb{P}(X - \mathbb{E}[X] \leq -t) = \mathbb{P}(-X - \mathbb{E}[-X] \geq t) \quad (1.26)$$

$$\leq \exp \left(-\frac{t^2}{2\sigma^2} \right), \quad (1.27)$$

where the last step follows from the upper-tail bound applied to $-X$. $\square$

Now if we combine upper and lower tails with the union bound yields the two-sided concentration inequality we will get the following

Lemma 1.3.5. *If X is sub-Gaussian with parameter $\sigma > 0$,*

$$\mathbb{P}(|X - \mathbb{E}X| \geq t) \leq 2 \exp \left(-\frac{t^2}{2\sigma^2} \right), \quad t > 0. \quad (1.28)$$

Note: This much faster decay is one reason sub-Gaussian assumptions are central in non-asymptotic high-dimensional statistics [1].

Example 1.3.6 (Rademacher random variable). A Rademacher random variable ε takes values in $\{-1, +1\}$ equiprobably:

In particular, $\mathbb{E}[\varepsilon] = 0$. Its MGF is

$$\mathbb{E}[e^{\lambda\varepsilon}] = \frac{1}{2}e^\lambda + \frac{1}{2}e^{-\lambda} \quad (1.29)$$

$$= \frac{1}{2} \left(\sum_{k=0}^{\infty} \frac{\lambda^k}{k!} + \sum_{k=0}^{\infty} \frac{(-\lambda)^k}{k!} \right) \quad (1.30)$$

$$= \sum_{m=0}^{\infty} \frac{\lambda^{2m}}{(2m)!}, \quad (1.31)$$

where the odd-order terms cancel.

Since $(2m)! \geq 2^m m!$,

$$\mathbb{E}[e^{\lambda\varepsilon}] \leq \sum_{m=0}^{\infty} \frac{\lambda^{2m}}{2^m m!} \quad (1.32)$$

$$= e^{\lambda^2/2}. \quad (1.33)$$

Since $\mathbb{E}[\varepsilon] = 0$, it follows that

$$\mathbb{E}[e^{\lambda(\varepsilon - \mathbb{E}[\varepsilon])}] \leq e^{\lambda^2/2}.$$

Therefore, ε is sub-Gaussian with parameter $\sigma = 1$.

This example shows that sub-Gaussianity is not the same as Gaussianity. A Rademacher variable is discrete and bounded, yet its MGF is dominated by that of a standard Gaussian. The sub-Gaussian class is therefore broad enough to include many non-Gaussian random variables while retaining Gaussian-like tail behavior.

Note: "equiprobably" meaning

$$\mathbb{P}(\varepsilon = 1) = \mathbb{P}(\varepsilon = -1) = \frac{1}{2}. \quad (1.34)$$

Therefore,

$$\mathbb{E}[\varepsilon] = \frac{1}{2}(1) + \frac{1}{2}(-1) = 0. \quad (1.35)$$

1.3.3 Example of Sub-Gaussian Variables

Example 1.3.7 (Bounded random variables). Let X be zero-mean, and supported on some interval $[a, b]$. We will show that X is sub-Gaussian with parameter at most $\sigma = b - a$.

Letting X' be an independent copy of X , so that $X \stackrel{d}{=} X'$ and $X \perp\!\!\!\perp X'$. For any $\lambda \in \mathbb{R}$, we have

$$\mathbb{E}_X [e^{\lambda X}] = \mathbb{E}_X [e^{\lambda(X - \mathbb{E}_{X'}[X'])}] \stackrel{(i)}{\leq} \mathbb{E}_X [\mathbb{E}_{X'} e^{\lambda(X - X')}] = \mathbb{E}_{X, X'} [e^{\lambda(X - X')}]$$

The first equality is because $\mathbb{E}[X'] = 0$, and (i) follows from the the convexity of the exponential, and Jensen's inequality. Letting ε be an independent Rademacher variable, note that since (X, X') are i.i.d., $X - X' \stackrel{d}{=} X' - X$, so $(X - X')$ is symmetric about 0. Thus, the distribution of $(X - X')$ is the same as that of $\varepsilon(X - X')$, so that we have

$$\mathbb{E}_{X, X'} [e^{\lambda(X - X')}] = \mathbb{E}_{X, X'} [\mathbb{E}_\varepsilon [e^{\lambda\varepsilon(X - X')}]] \stackrel{(ii)}{\leq} \mathbb{E}_{X, X'} [e^{\frac{\lambda^2}{2}(X - X')^2}] \leq e^{\frac{\lambda^2}{2}(b-a)^2}$$

The first equality follows from Tonelli's theorem and (ii) follows from the sub-Gaussianity of Radamacher variable and last inequality is because $|X - X'| \leq b - a$. Putting everything together, we have shown that X is sub-Gaussian with parameter at most $\sigma = b - a$.

Remark 1.3.8. This result is useful, but it can be sharpened. It can be shown that X is sub-Gaussian with parameter at most $\sigma = \frac{b-a}{2}$ (Exercise 2.4). The technique used in this example is a simple example of a *symmetrization argument*, in which we first introduce an independent copy X' , and then symmetrize the problem with a Rademacher variable. Such symmetrization arguments are useful in a variety of contexts.

1.3.4 Hoeffding's Inequality

Another property of sub-Gaussianity is preserved by linear operations.

Lemma 1.3.9. *Let X_1 and X_2 be independent sub-Gaussian variables with parameters σ_1 and σ_2 , respectively, then $X_1 + X_2$ is sub-Gaussian with parameter $\sqrt{\sigma_1^2 + \sigma_2^2}$.*

Proof. By independence we have

$$\mathbb{E} [e^{\lambda(X_1 + X_2 - \mathbb{E}X_1 - \mathbb{E}X_2)}] = \mathbb{E} [e^{\lambda(X_1 - \mathbb{E}X_1)}] \mathbb{E} [e^{\lambda(X_2 - \mathbb{E}X_2)}] \leq e^{\frac{\lambda^2}{2}\sigma_1^2} e^{\frac{\lambda^2}{2}\sigma_2^2} = e^{\frac{\lambda^2}{2}(\sigma_1^2 + \sigma_2^2)}$$

□

As a consequence of this fact, we obtain an important result known as the *Hoeffding's inequality*:

Proposition 1.3.10 (Hoeffding's inequality). *Suppose that the variables X_i , $i = 1, \dots, n$, are independent, and X_i has mean μ_i and sub-Gaussian parameter σ_i . Then for all $t \geq 0$, we have*

$$\mathbb{P} \left[\sum_{i=1}^n (X_i - \mu_i) \geq t \right] \leq \exp \left\{ -\frac{t^2}{2 \sum_{i=1}^n \sigma_i^2} \right\}$$

and

$$\mathbb{P} \left[\left| \sum_{i=1}^n X_i - \sum_{i=1}^n \mathbb{E} X_i \right| \geq t \right] \leq 2 \exp \left\{ -\frac{t^2}{2 \sum_{i=1}^n \sigma_i^2} \right\}$$

Proof. By previous lemma, $\sum_{i=1}^n X_i$ is $\sqrt{\sigma_1^2 + \dots + \sigma_n^2}$ sub-Gaussian. So, by basic sub-Gaussian tail bound we have the results. $\square$

Corollary 1.3.11. *If X_i , $i = 1, \dots, n$, are i.i.d, with mean μ and sub-Gaussian parameter σ . Then for all $t \geq 0$, we have*

$$\mathbb{P} \left[\left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right| \geq t \right] \leq 2 \exp \left\{ -\frac{nt^2}{2\sigma^2} \right\}$$

and

$$\mathbb{P} \left[\sqrt{n} \left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right| \geq t \right] \leq 2 \exp \left\{ -\frac{t^2}{2\sigma^2} \right\}$$

Remark 1.3.12. The above can be viewed as a finite sample Central Limit Theorem. CLT states that when $n \rightarrow \infty$,

$$\sqrt{n} \left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right| \rightarrow \mathcal{N}(0, \sigma^2).$$

The second statement of the above corollary is a finite sample version of CLT, by stating that the tail of $\sqrt{n} \left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right|$ is upper bounded by the tail of $\mathcal{N}(0, \sigma^2)$.

Proof. By Hoeffding,

$$\mathbb{P} \left[\left| \frac{1}{n} \sum_{i=1}^n X_i - \mu \right| \geq t \right] = \mathbb{P} \left[\left| \sum_{i=1}^n X_i - n\mu \right| \geq nt \right] \leq 2 \exp \left\{ -\frac{(nt)^2}{2n\sigma^2} \right\} = 2 \exp \left\{ -\frac{nt^2}{2\sigma^2} \right\}$$

Second statement follows analogously. $\square$

One application of Hoeffding's bound is to derive finite sample guarantees for the error of the sample mean.

Example 1.3.13 (Finite-sample error guarantee). Let $\bar{X}_n := \frac{1}{n} \sum_{i=1}^n X_i$ and fix a confidence level $\delta \in (0, 1)$. By the preceding corollary,

$$\mathbb{P} [|\bar{X}_n - \mu| \geq t] \leq 2 \exp \left\{ -\frac{nt^2}{2\sigma^2} \right\}$$

Choose $t = \sqrt{\frac{2\sigma^2}{n} \log \left(\frac{2}{\delta} \right)}$, so that $\delta = 2 \exp \left\{ -\frac{nt^2}{2\sigma^2} \right\}$. In other words, with probability at least $1 - \delta$,

$$|\bar{X}_n - \mu| \leq \sqrt{\frac{2\sigma^2}{n} \log \left(\frac{2}{\delta} \right)}$$

Example 1.3.14 (Expected absolute error). We can also control the expected absolute error $\mathbb{E}|\bar{X}_n - \mu|$ directly from the tail bound. Recall that for any nonnegative random variable Y , Tonelli's theorem gives

$$\mathbb{E}[Y] = \int_0^\infty \mathbb{P}(Y \geq t) dt$$

Taking $Y = |\bar{X}_n - \mu|$, the preceding corollary and the trivial bound $\mathbb{P}(Y \geq t) \leq 1$ imply

$$\mathbb{P} [|\bar{X}_n - \mu| \geq t] \leq \min \left\{ 1, 2 \exp \left\{ -\frac{nt^2}{2\sigma^2} \right\} \right\}$$

Let t_0 be the point at which the two upper bounds agree. Thus $2e^{-\frac{nt_0^2}{2\sigma^2}} = 1$, which gives $t_0 = \sigma \sqrt{\frac{2 \log 2}{n}}$. Hence

$$\begin{aligned} \mathbb{E} |\bar{X}_n - \mu| &= \int_0^\infty \mathbb{P} [|\bar{X}_n - \mu| \geq t] dt \\ &\leq \int_0^{t_0} 1 dt + \underbrace{\int_{t_0}^\infty 2 \exp \left\{ -\frac{nt^2}{2\sigma^2} \right\} dt}_{z = \frac{\sqrt{n}}{\sigma} t, \quad dt = \frac{\sigma}{\sqrt{n}} dz} \\ &= t_0 + \frac{2\sigma}{\sqrt{n}} \int_{\frac{\sqrt{n}}{\sigma} t_0}^\infty e^{-z^2/2} dz \\ &= \frac{\sigma}{\sqrt{n}} \left(\sqrt{2 \log 2} + 2 \int_{\sqrt{2 \log 2}}^\infty e^{-z^2/2} dz \right) \\ &\leq 2 \frac{\sigma}{\sqrt{n}}. \end{aligned}$$

We conclude our discussion of sub-Gaussianity with a theorem that provides different characterizations of sub-Gaussian variables.

Theorem 1.3.15 (Equivalent characterizations of sub-Gaussian variables). *Given any zero-mean random variable X , the following properties are equivalent:*

1. *There is a constant $\sigma \geq 0$ such that*

$$\mathbb{E} [e^{\lambda X}] \leq e^{\frac{\lambda^2 \sigma^2}{2}} \quad \text{for all } \lambda \in \mathbb{R}.$$

2. *There is a constant $c \geq 0$ and Gaussian random variable $Z \sim \mathcal{N}(0, \tau^2)$ such that*

$$\mathbb{P} [|X| \geq s] \leq c \mathbb{P} [|Z| \geq s] \quad \text{for all } s \geq 0.$$

3. *There is a constant $\theta \geq 0$ such that*

$$\mathbb{E} [X^{2k}] \leq \frac{(2k)!}{2^k k!} \theta^{2k} \quad \text{for all } k = 1, 2, \dots$$

4. *There is a constant $\sigma \geq 0$ such that*

$$\mathbb{E} \left[e^{\frac{\lambda X^2}{2\sigma^2}} \right] \leq \frac{1}{\sqrt{1-\lambda}} \quad \text{for all } \lambda \in [0, 1).$$

Proof. See Appendix A (Section 2.4) in [1]. □

1.4 Symmetrization

Lemma 1.4.1. *Let $Y \perp\!\!\!\perp Z$ with $\mathbb{E}[Z] = 0$. Then for any convex function F ,*

$$\mathbb{E}[F(Y)] \leq \mathbb{E}[F(Y + Z)]. \quad (1.36)$$

Example 1.4.2. The function $F(x) = e^{\lambda x}$ is convex. Suppose $\mathbb{E}[X] = 0$. Applying Lemma 1.4.1 with $Y = X$ and $Z = -X'$, where $X' \perp\!\!\!\perp X$ and $X' \sim \text{Law}(X)$, we obtain

$$\mathbb{E}_X [e^{\lambda X}] \leq \mathbb{E}_{X, X'} [e^{\lambda(X - X')}] \quad (1.37)$$

Proof of Lemma 1.4.1. For every $y \in \mathbb{R}$, using $\mathbb{E}[Z] = 0$ and Jensen's inequality,

$$F(y) = F(y + \mathbb{E}_Z[Z]) = F(\mathbb{E}_Z[y + Z]) \leq \mathbb{E}_Z[F(y + Z)].$$

Substituting $y = Y$ gives $F(Y) \leq \mathbb{E}_Z[F(Y + Z)]$, and taking expectations over Y yields

$$\mathbb{E}_Y[F(Y)] \leq \mathbb{E}_Y \mathbb{E}_Z[F(Y + Z)] = \mathbb{E}_{Y, Z}[F(Y + Z)],$$

where the last equality uses $Y \perp\!\!\!\perp Z$. □

Definition 1.4.3. A random variable X is called *symmetric* if it has the same distribution as $-X$.

Remark 1.4.4. If X has pdf $f(x)$, then X is symmetric if and only if $f(x) = f(-x)$.

Lemma 1.4.5 (Constructing symmetric distributions). *Let X be a random variable and let ε be an independent Rademacher random variable.*

- (a) εX and $\varepsilon |X|$ are identically distributed and symmetric.
- (b) If X is symmetric, then both εX and $\varepsilon |X|$ have the same distribution as X .
- (c) If X' is an independent copy of X , then $X - X'$ is symmetric.

Proof. We only prove that εX is symmetric. For every event $A \subseteq \mathbb{R}$, conditioning on the value of ε gives

$$\mathbb{P}[\varepsilon X \in A] = \frac{1}{2} \mathbb{P}[X \in A] + \frac{1}{2} \mathbb{P}[-X \in A] = \frac{1}{2} \mathbb{P}[-X \in A] + \frac{1}{2} \mathbb{P}[X \in A] = \mathbb{P}[-\varepsilon X \in A]. \quad \square$$

Lemma 1.4.6 (Symmetrization). *Let $X_1, \dots, X_N$ be independent, mean-zero random vectors in some normed space, and let $\varepsilon_1, \dots, \varepsilon_N$ be independent Rademacher random variables. Then*

$$\frac{1}{2} \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X_i \right\| \leq \mathbb{E} \left\| \sum_{i=1}^N X_i \right\| \leq 2 \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X_i \right\|. \quad (1.38)$$

Proof. Let $(X'_1, \dots, X'_N)$ be independent copies of $(X_1, \dots, X_N)$. By Lemma 1.4.5(c) the vectors $X_i - X'_i$ are independent and symmetric (this step uses the assumption that X_i are independent), so that

$$(X_1 - X'_1, \dots, X_N - X'_N) \stackrel{d}{=} (\varepsilon_1(X_1 - X'_1), \dots, \varepsilon_N(X_N - X'_N)). \quad (1.39)$$

Upper bound. Set $p := \mathbb{E} \left\| \sum_{i=1}^N X_i \right\|$. Since $\mathbb{E} \left[\sum_{i=1}^N X'_i \right] = 0$, we may apply Lemma 1.4.1 with $F(x) = \|x\|$, $Y = \sum_{i=1}^N X_i$ and $Z = -\sum_{i=1}^N X'_i$ to get

$$p = \mathbb{E} \left\| \sum_{i=1}^N X_i - \mathbb{E} \left[\sum_{i=1}^N X'_i \right] \right\| \leq \mathbb{E} \left\| \sum_{i=1}^N (X_i - X'_i) \right\|.$$

Then by (1.39) and the triangle inequality,

$$\begin{aligned} \mathbb{E} \left\| \sum_{i=1}^N (X_i - X'_i) \right\| &= \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i (X_i - X'_i) \right\| = \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X_i - \sum_{i=1}^N \varepsilon_i X'_i \right\| \\ &\leq \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X_i \right\| + \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X'_i \right\| = 2 \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X_i \right\|. \end{aligned}$$

Lower bound. Conditionally on $\varepsilon = (\varepsilon_1, \dots, \varepsilon_N)$, the vectors $\sum_{i=1}^N \varepsilon_i X_i$ and $-\sum_{i=1}^N \varepsilon_i X'_i$ are independent and the latter has mean zero, so Lemma 1.4.1 applies conditionally and gives, after averaging over ε ,

$$\mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X_i \right\| \leq \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i X_i - \sum_{i=1}^N \varepsilon_i X'_i \right\| = \mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i (X_i - X'_i) \right\|.$$

Applying (1.39) once more and then the triangle inequality,

$$\mathbb{E} \left\| \sum_{i=1}^N \varepsilon_i (X_i - X'_i) \right\| = \mathbb{E} \left\| \sum_{i=1}^N (X_i - X'_i) \right\| \leq \mathbb{E} \left\| \sum_{i=1}^N X_i \right\| + \mathbb{E} \left\| \sum_{i=1}^N X'_i \right\| = 2 \mathbb{E} \left\| \sum_{i=1}^N X_i \right\|. \quad \square$$

Remark 1.4.7. The zero-mean assumption is needed for both bounds: in the upper bound to subtract $\mathbb{E}[\sum_i X'_i]$, and in the lower bound to guarantee that $\mathbb{E}[\sum_i \varepsilon_i X'_i \mid \varepsilon] = 0$.

1.5 Sub-exponential Random Variables

We now turn to the class of sub-exponential variables, which are defined by a slightly milder condition on the moment generating function.

Definition 1.5.1 ([1, Definition 2.7]). A random variable X with mean $\mu = \mathbb{E}[X]$ is *sub-exponential* if there are non-negative parameters (ν, α) such that

$$\mathbb{E} \left[e^{\lambda(X-\mu)} \right] \leq e^{\frac{\nu^2 \lambda^2}{2}} \quad \text{for all } |\lambda| < \frac{1}{\alpha}. \quad (1.40)$$

Remark 1.5.2. A σ -sub-Gaussian variable is $(\sigma, 0)$ -sub-exponential, where we interpret $1/0$ as $+\infty$.

Example 1.5.3 (Sub-exponential but not sub-Gaussian). Let $Z \sim N(0, 1)$ and consider $X = Z^2$, so that $\mathbb{E}[X] = 1$. For $\lambda < \frac{1}{2}$,

$$\begin{aligned} \mathbb{E} \left[e^{\lambda(X-1)} \right] &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{+\infty} e^{\lambda(z^2-1)} e^{-\frac{z^2}{2}} dz \\ &= \frac{e^{-\lambda}}{\sqrt{2\pi}} \int_{-\infty}^{+\infty} e^{-\frac{(1-2\lambda)z^2}{2}} dz = \frac{e^{-\lambda}}{\sqrt{1-2\lambda}}, \end{aligned} \quad (1.41)$$

where the last equality uses $\int_{-\infty}^{+\infty} e^{-az^2/2} dz = \sqrt{2\pi/a}$ with $a = 1 - 2\lambda > 0$. For $\lambda \geq \frac{1}{2}$ the moment generating function is infinite, so X is not sub-Gaussian. On the other hand,

$$\frac{e^{-\lambda}}{\sqrt{1-2\lambda}} \leq e^{2\lambda^2} = e^{\frac{2^2 \lambda^2}{2}} \quad \text{where } |\lambda| < \frac{1}{4} \quad (\text{check this!}), \quad (1.42)$$

so that $\mathbb{E} \left[e^{\lambda(X-1)} \right] \leq e^{\frac{2^2 \lambda^2}{2}}$ for all $|\lambda| < \frac{1}{4}$, i.e. X is $(2, 4)$ -sub-exponential.

Exercises

1. Verify the inequality (1.42), that is, show that $\frac{e^{-\lambda}}{\sqrt{1-2\lambda}} \leq e^{2\lambda^2}$ for all $|\lambda| < \frac{1}{4}$.

Proposition 1.5.4 (Sub-exponential tail bound). *Suppose X is sub-exponential with parameters (ν, α) and mean μ . Then*

$$\mathbb{P}[X - \mu \geq t] \leq \begin{cases} e^{-\frac{t^2}{2\nu^2}}, & \text{if } 0 \leq t \leq \frac{\nu^2}{\alpha}, \\ e^{-\frac{t}{2\alpha}}, & \text{for } t > \frac{\nu^2}{\alpha}, \end{cases} = e^{-\min\left\{\frac{t^2}{2\nu^2}, \frac{t}{2\alpha}\right\}}. \quad (1.43)$$

Proof. Assume $\mu = 0$ without loss of generality. By the Chernoff bound from Lecture 1 and Definition 1.5.1,

$$\begin{aligned} \log \mathbb{P}[X \geq t] &\leq \inf_{0 \leq \lambda < \frac{1}{\alpha}} \left\{ \log \mathbb{E}[e^{\lambda X}] - \lambda t \right\} \leq \inf_{0 \leq \lambda < \frac{1}{\alpha}} \left\{ \log e^{\frac{\nu^2 \lambda^2}{2}} - \lambda t \right\} \\ &= \inf_{0 \leq \lambda < \frac{1}{\alpha}} \underbrace{\left\{ \frac{\nu^2 \lambda^2}{2} - \lambda t \right\}}_{=: g(\lambda)}. \end{aligned} \quad (1.44)$$

The function g is a convex parabola (see Figure 1.2) with

$$g'(\lambda) = \nu^2 \lambda - t = 0 \quad \Longleftrightarrow \quad \lambda = \frac{t}{\nu^2}.$$

- ① If $\frac{t}{\nu^2} < \frac{1}{\alpha}$ (i.e. $t < \frac{\nu^2}{\alpha}$), the unconstrained minimizer lies in the feasible region, so

$$\inf_{0 \leq \lambda < \frac{1}{\alpha}} g(\lambda) = g\left(\frac{t}{\nu^2}\right) = \frac{\nu^2}{2} \cdot \frac{t^2}{\nu^4} - \frac{t}{\nu^2} \cdot t = -\frac{t^2}{2\nu^2}.$$

- ② If $\frac{t}{\nu^2} \geq \frac{1}{\alpha}$ (i.e. $t \geq \frac{\nu^2}{\alpha}$), then g is decreasing on $[0, \frac{t}{\nu^2})$ and the infimum is approached at $\lambda = \frac{1}{\alpha}$, giving

$$\inf_{0 \leq \lambda < \frac{1}{\alpha}} g(\lambda) = g\left(\frac{1}{\alpha}\right) = \frac{\nu^2}{2} \cdot \frac{1}{\alpha^2} - \frac{1}{\alpha} t = \frac{\nu^2}{2\alpha^2} - \frac{t}{\alpha} \leq -\frac{t}{2\alpha},$$

where the last step uses $\nu^2 < \alpha t$. □

1.5.1 Bernstein's Condition

We say that X satisfies *Bernstein's condition with parameter b* if

$$|\mathbb{E}[(X - \mu)^k]| \leq \frac{1}{2} k! \operatorname{Var}(X) b^{k-2}, \quad \text{for any } k = 2, 3, \dots, \quad (1.45)$$

where $\mu = \mathbb{E}[X]$.

Example 1.5.5. If (1.45) is satisfied, then X is sub-exponential. Indeed, writing $\sigma^2 = \operatorname{Var}(X)$ and using the Taylor expansion of the exponential together with $\mathbb{E}[X - \mu] = 0$,

$$\begin{aligned} \mathbb{E}[e^{\lambda(X-\mu)}] &= \mathbb{E}\left[\sum_{k=0}^{\infty} \frac{(\lambda(X-\mu))^k}{k!}\right] = 1 + 0 + \mathbb{E}\left[\sum_{k=2}^{\infty} \frac{\lambda^k (X-\mu)^k}{k!}\right] \\ &= 1 + \sum_{k=2}^{\infty} \frac{\lambda^k \mathbb{E}[(X-\mu)^k]}{k!} = 1 + \frac{\lambda^2 \sigma^2}{2} + \sum_{k=3}^{\infty} \frac{\lambda^k \mathbb{E}[(X-\mu)^k]}{k!} \\ &\leq 1 + \frac{\lambda^2 \sigma^2}{2} + \sum_{k=3}^{\infty} \frac{|\lambda|^k \cdot \frac{1}{2} k! \sigma^2 b^{k-2}}{k!} = 1 + \frac{\lambda^2 \sigma^2}{2} + \frac{\lambda^2 \sigma^2}{2} \sum_{k=3}^{\infty} (|\lambda| b)^{k-2} = 1 + \frac{\lambda^2 \sigma^2}{2} \sum_{s=0}^{\infty} (|\lambda| b)^s. \end{aligned} \quad (1.46)$$




Figure 1.2: The function $g(\lambda) = \frac{\nu^2 \lambda^2}{2} - \lambda t$ appearing in (1.44), in the case $\frac{t}{\nu^2} > \frac{1}{\alpha}$. The open circle marks the unconstrained minimizer $\frac{t}{\nu^2}$, which lies outside the feasible region $|\lambda| < \frac{1}{\alpha}$.

For $|\lambda| < \frac{1}{b}$ the geometric series sums to $\frac{1}{1-|\lambda|b}$, so that

$$\mathbb{E} \left[e^{\lambda(X-\mu)} \right] \leq 1 + \frac{\lambda^2 \sigma^2 / 2}{1 - b|\lambda|} \leq e^{\frac{\lambda^2 \sigma^2 / 2}{1 - b|\lambda|}}, \quad (1.47)$$

using $1 + t \leq e^t$. Restricting further to $|\lambda| < \frac{1}{2b}$ gives $1 - b|\lambda| > \frac{1}{2}$, and therefore

$$\mathbb{E} \left[e^{\lambda(X-\mu)} \right] \leq e^{\lambda^2 \sigma^2} = e^{\frac{\lambda^2 (\sqrt{2}\sigma)^2}{2}}, \quad \text{for all } |\lambda| < \frac{1}{2b},$$

i.e. X is sub-exponential with parameters $(\nu, \alpha) = (\sqrt{2}\sigma, 2b)$.

Proposition 1.5.6 (Bernstein-type bound). *Suppose X satisfies Bernstein's condition (1.45) with parameter b , mean μ , and variance σ^2 . Then*

$$\mathbb{E} \left[e^{\lambda(X-\mu)} \right] \leq \exp \left(\frac{\lambda^2 \sigma^2 / 2}{1 - b|\lambda|} \right) \quad \text{for all } |\lambda| < \frac{1}{b}, \quad (1.48)$$

and, moreover, the concentration inequality

$$\mathbb{P}(|X - \mu| > t) \leq 2 \exp \left(-\frac{t^2}{2(\sigma^2 + bt)} \right) \quad \text{for all } t \geq 0. \quad (1.49)$$

Proof. Equation (1.48) directly follows from (1.46).

Remark 1.5.7. Restricting further to $|\lambda| < \frac{1}{2b}$ gives $b|\lambda| < \frac{1}{2}$, hence $1 - b|\lambda| > \frac{1}{2}$ and

$$\mathbb{E} \left[e^{\lambda X} \right] \leq \exp \left(\frac{\lambda^2 \sigma^2 / 2}{1/2} \right) = e^{\lambda^2 \sigma^2} = e^{\frac{\lambda^2 (\sqrt{2}\sigma)^2}{2}}, \quad |\lambda| < \frac{1}{2b},$$

which recovers the Topic 3 conclusion that X is $(\sqrt{2}\sigma, 2b)$ -sub-exponential. Bound (1.48) is sharper: it holds on the full range $|\lambda| < 1/b$ instead of only $|\lambda| < 1/(2b)$, and this extra room is exactly what lets us optimize the Chernoff bound below.

Proof of (1.49). By the Chernoff bound and (1.48), for any $t \geq 0$,

$$\log \mathbb{P}(X - \mu \geq t) \leq \inf_{0 \leq \lambda < 1/b} \left\{ \log \mathbb{E} \left[e^{\lambda(X - \mu)} \right] - \lambda t \right\} \leq \inf_{0 \leq \lambda < 1/b} \underbrace{\left\{ \frac{\lambda^2 \sigma^2}{2(1 - b\lambda)} - \lambda t \right\}}_{=: \theta(\lambda)}. \quad (1.50)$$

Rather than optimizing θ exactly, it suffices to evaluate it at the specific point

$$\lambda^* := \frac{t}{\sigma^2 + bt}.$$

Since $t, \sigma^2, b \geq 0$ we have $\lambda^* \geq 0$, and $\lambda^* < 1/b \iff bt < \sigma^2 + bt \iff \sigma^2 > 0$, which holds whenever X is non-degenerate (if $\sigma^2 = 0$ then $X \equiv \mu$ a.s. and (1.49) is trivial). At $\lambda = \lambda^*$,

$$1 - b\lambda^* = 1 - \frac{bt}{\sigma^2 + bt} = \frac{\sigma^2}{\sigma^2 + bt},$$

so that

$$\theta(\lambda^*) = \frac{\sigma^2(\lambda^*)^2(\sigma^2 + bt)}{2\sigma^2} - \lambda^* t = \frac{(\lambda^*)^2(\sigma^2 + bt)}{2} - \lambda^* t = \frac{t^2}{2(\sigma^2 + bt)} - \frac{t^2}{\sigma^2 + bt} = -\frac{t^2}{2(\sigma^2 + bt)}.$$

Combining with (1.50) (using that λ^* is feasible),

$$\mathbb{P}(X - \mu \geq t) \leq \exp \left(-\frac{t^2}{2(\sigma^2 + bt)} \right), \quad t \geq 0. \quad (1.51)$$

Finally, $-X$ also satisfies Bernstein's condition with the same (σ^2, b) , since $|\mathbb{E}[(-(X - \mu))^k]| = |\mathbb{E}[(X - \mu)^k]|$ and $\text{Var}(-X) = \sigma^2$. Applying (1.51) to $-X$ gives $\mathbb{P}(X - \mu \leq -t) \leq \exp(-t^2/(2(\sigma^2 + bt)))$ as well, and adding the two one-sided bounds yields (1.49). $\square$

Exercise 1.5.8. Compare the bound (1.49) with the bound one would obtain by plugging the (weaker) parameters $(\sqrt{2}\sigma, 2b)$ derived in the Remark above directly into the general sub-exponential tail bound of Topic 3. Show that (1.49) is never worse, and can be substantially tighter when $bt \ll \sigma^2$.

1.5.2 Sums of Independent Sub-exponential Random Variables

Recall that X with mean μ is *sub-exponential with parameters* (ν, α) if $\mathbb{E}[e^{\lambda(X - \mu)}] \leq e^{\nu^2 \lambda^2 / 2}$ for all $|\lambda| < 1/\alpha$, and that in this case

$$\mathbb{P}(|X - \mu| > t) \leq \begin{cases} 2e^{-t^2/(2\nu^2)}, & 0 \leq t \leq \nu^2/\alpha, \\ 2e^{-t/(2\alpha)}, & t > \nu^2/\alpha. \end{cases} \quad (1.52)$$

Lemma 1.5.9. Let $X_1, \dots, X_n$ be independent random variables, with X_i sub-exponential with parameters (ν_i, d_i) . Then $S_n := \sum_{i=1}^n X_i$ is sub-exponential with parameters

$$\left(\sqrt{\sum_{i=1}^n \nu_i^2}, \max_{1 \leq i \leq n} d_i \right).$$

and

$$\mathbb{P}(|S_n - \mathbb{E}[S_n]| > t) \leq \begin{cases} 2e^{-t^2/(2V^2)}, & 0 \leq t \leq V^2/d, \\ 2e^{-t/(2d)}, & t > V^2/d, \end{cases} \quad V^2 = \sum_{i=1}^n \nu_i^2, \quad d = \max_{1 \leq i \leq n} d_i. \quad (1.53)$$

Proof. Write $d := \max_{1 \leq i \leq n} d_i$ and $V^2 := \sum_{i=1}^n \nu_i^2$. For any $|\lambda| < 1/d = \min_{1 \leq i \leq n} (1/d_i)$, we have $|\lambda| < 1/d_i$ for every i , so by independence,

$$\mathbb{E} \left[e^{\lambda(S_n - \mathbb{E}[S_n])} \right] = \mathbb{E} \left[e^{\lambda \sum_{i=1}^n (X_i - \mathbb{E}[X_i])} \right] = \prod_{i=1}^n \mathbb{E} \left[e^{\lambda(X_i - \mathbb{E}[X_i])} \right] \leq \prod_{i=1}^n e^{\nu_i^2 \lambda^2 / 2} = e^{\lambda^2 V^2 / 2}.$$

Hence S_n is (V, d) -sub-exponential.

By (1.52) applied with $(\nu, \alpha) = (V, d)$, we obtain the tail bound

$$\mathbb{P}(|S_n - \mathbb{E}[S_n]| > t) \leq \begin{cases} 2e^{-t^2/(2V^2)}, & 0 \leq t \leq V^2/d, \\ 2e^{-t/(2d)}, & t > V^2/d, \end{cases} \quad V^2 = \sum_{i=1}^n \nu_i^2, \quad d = \max_{1 \leq i \leq n} d_i.$$

□

Example 1.5.10 (Concentration of χ^2 variables). Let $Z_1, \dots, Z_n$ i.i.d. $N(0, 1)$ and let $Y := \sum_{i=1}^n Z_i^2 \sim \chi_n^2$, so that $\mathbb{E}[Y] = n$. Recall from Topic 3 that each Z_i^2 has mean 1 and is $(2, 4)$ -sub-exponential, i.e. $\mathbb{E}[e^{\lambda(Z_i^2 - 1)}] \leq e^{2\lambda^2}$ for all $|\lambda| < 1/4$. Applying Lemma 1.5.9 with $\nu_i = 2$, $d_i = 4$ for every i , Y is $(\sqrt{4n}, 4) = (2\sqrt{n}, 4)$ -sub-exponential; here $V^2 = 4n$ and $d = 4$, so the threshold in (1.53) is $V^2/d = n$. Thus

$$\mathbb{P}(|Y - n| \geq s) \leq \begin{cases} 2e^{-s^2/(8n)}, & 0 \leq s \leq n, \\ 2e^{-s/8}, & s > n. \end{cases}$$

Substituting $s = nt$ (so $s \leq n \iff t \leq 1$) gives

$$\mathbb{P}\left(\left|\frac{Y}{n} - 1\right| \geq t\right) \leq \begin{cases} 2e^{-nt^2/8}, & 0 \leq t \leq 1, \\ 2e^{-nt/8}, & t > 1. \end{cases} \quad (1.54)$$

1.5.3 Application: The Johnson–Lindenstrauss Lemma

Fix distinct points $u^1, \dots, u^N \in \mathbb{R}^d$, where d is large. We would like a map $F : \mathbb{R}^d \rightarrow \mathbb{R}^m$ with $m \ll d$, so that we only need to store the compressed points $\{F(u^1), \dots, F(u^N)\}$, while approximately preserving all pairwise (squared) distances: given a distortion parameter $\delta \in (0, 1)$, we want

$$1 - \delta \leq \frac{\|F(u^i) - F(u^j)\|_2^2}{\|u^i - u^j\|_2^2} \leq 1 + \delta, \quad \text{for all } i \neq j. \quad (1.55)$$

Construction. Let $X \in \mathbb{R}^{m \times d}$ have i.i.d. $N(0, 1)$ entries, and write its rows as $X_1, \dots, X_m \in \mathbb{R}^d$, so that $X_1, \dots, X_m$ i.i.d. $N(0, I_d)$. Define the (random) linear map

$$F(u) := \frac{Xu}{\sqrt{m}}, \quad u \in \mathbb{R}^d.$$

Step 1: distribution of $\|Xu\|_2^2$. Fix $u \in \mathbb{R}^d$, $u \neq 0$. Writing $Xu = (\langle X_1, u \rangle, \dots, \langle X_m, u \rangle)^\top$, for each i and any unit vector $v = u/\|u\|_2$,

$$\langle X_i, v \rangle = \sum_{k=1}^d X_{i,k} v_k \sim N\left(0, \sum_{k=1}^d v_k^2\right) = N(0, 1) \quad (\text{check: } \|v\|_2 = 1), \quad (1.56)$$

since $X_i \sim N(0, I_d)$ and any linear combination of independent standard normals with coefficient vector of unit ℓ_2 -norm is again standard normal. Moreover $\langle X_1, v \rangle, \dots, \langle X_m, v \rangle$ are independent because $X_1, \dots, X_m$ are independent. Hence

$$Y := \frac{\|Xu\|_2^2}{\|u\|_2^2} = \sum_{i=1}^m \frac{\langle X_i, u \rangle^2}{\|u\|_2^2} = \sum_{i=1}^m \langle X_i, v \rangle^2 \sim \chi_m^2. \quad (1.57)$$

Step 2: concentration for a single vector. Since $\|F(u)\|_2^2 / \|u\|_2^2 = \|Xu\|_2^2 / (m \|u\|_2^2) = Y/m$, applying (1.54) (with n replaced by m) gives, for any $\delta \in (0, 1)$,

$$\mathbb{P} \left(\frac{\|F(u)\|_2^2}{\|u\|_2^2} \notin (1 - \delta, 1 + \delta) \right) = \mathbb{P} \left(\left| \frac{Y}{m} - 1 \right| \geq \delta \right) \leq 2e^{-m\delta^2/8}. \quad (1.58)$$

Step 3: union bound over all pairs. Apply (1.58) with $u = u^i - u^j$ for each of the $\binom{N}{2}$ unordered pairs $i \neq j$ (assuming $u^i \neq u^j$):

$$\mathbb{P} \left(\frac{\|F(u^i) - F(u^j)\|_2^2}{\|u^i - u^j\|_2^2} \notin (1 - \delta, 1 + \delta) \right) \leq 2e^{-m\delta^2/8}.$$

Let

$$I := \bigcup_{i < j} \left\{ \frac{\|F(u^i) - F(u^j)\|_2^2}{\|u^i - u^j\|_2^2} \notin (1 - \delta, 1 + \delta) \right\}$$

be the event that (1.55) fails for at least one pair. By the union bound,

$$\mathbb{P}(I) \leq \sum_{i < j} \mathbb{P} \left(\frac{\|F(u^i) - F(u^j)\|_2^2}{\|u^i - u^j\|_2^2} \notin (1 - \delta, 1 + \delta) \right) \leq \binom{N}{2} \cdot 2e^{-m\delta^2/8} = N(N-1)e^{-m\delta^2/8}. \quad (1.59)$$

Step 4: choosing m . Given a target failure probability $\varepsilon \in (0, 1)$, we want $\mathbb{P}(I) < \varepsilon$. By (1.59), it suffices that

$$N(N-1)e^{-m\delta^2/8} < \varepsilon \iff m > \frac{8}{\delta^2} \log \frac{N(N-1)}{\varepsilon}.$$

Since $N(N-1) \leq N^2$, a cleaner sufficient condition is

$$m > \frac{16}{\delta^2} \log \frac{N}{\varepsilon}. \quad (1.60)$$

Theorem 1.5.11 (Johnson–Lindenstrauss Lemma). *Let distinct $u^1, \dots, u^N \in \mathbb{R}^d$ and let $\delta, \varepsilon \in (0, 1)$. If*

$$m > \frac{16}{\delta^2} \log \frac{N}{\varepsilon},$$

then, with $X \in \mathbb{R}^{m \times d}$ having i.i.d. $N(0, 1)$ entries and $F(u) := Xu/\sqrt{m}$, with probability at least $1 - \varepsilon$ we have

$$1 - \delta \leq \frac{\|F(u^i) - F(u^j)\|_2^2}{\|u^i - u^j\|_2^2} \leq 1 + \delta \quad \text{simultaneously for all } i \neq j.$$

Proof. Immediate from (1.59) and (1.60): the choice of m forces $\mathbb{P}(I) < \varepsilon$, i.e. (1.55) holds for every pair on the complementary event, which has probability at least $1 - \varepsilon$. $\square$

Theorem 1.5.12 (Equivalent characterizations of sub-exponential random variables). *For a mean-zero random variable X , the following conditions are equivalent.*

(i) *There exist non-negative parameters (ν, α) such that*

$$\mathbb{E} [e^{\lambda X}] \leq e^{\nu^2 \lambda^2 / 2} \quad \text{for all } |\lambda| \leq \frac{1}{\alpha}.$$

(ii) *There exist constants $c_0, c_0 > 0$ such that*

$$\mathbb{E} [e^{\lambda X}] \leq c_0 \quad \text{for all } |\lambda| \leq c_0.$$

(iii) There exist constants $c_1, c_2 > 0$ such that

$$\mathbb{P}(|X| \geq t) \leq c_1 e^{-c_2 t} \quad \text{for all } t \geq 0.$$

(iv) The quantity

$$\gamma := \sup_{k \geq 2} \left(\frac{\mathbb{E}[|X|^k]}{k!} \right)^{1/k}$$

is finite.

Remark 1.5.13. A sub-exponential tail bound decays more slowly than a sub-Gaussian tail bound for large t , so the sub-exponential class permits heavier tails.

1.6 Martingale-Based Methods

Definition 1.6.1 (Martingale difference sequence). An adapted sequence $\{(D_k, \mathcal{F}_k)\}_{k=0}^\infty$ is called a *martingale difference sequence* if, for every $k \geq 1$,

$$\mathbb{E}[|D_k|] < \infty \quad \text{and} \quad \mathbb{E}[D_k \mid \mathcal{F}_{k-1}] = 0.$$

Equivalently, after shifting the index, $\mathbb{E}[D_{k+1} \mid \mathcal{F}_k] = 0$.

Remark 1.6.2. Here, D_k denotes the difference and $\mathcal{F}_k$ denotes the associated sigma-algebra.

Example 1.6.3. Let $\{(Y_k, \mathcal{F}_k)\}_{k \geq 0}$ be a martingale and define

$$D_k := Y_k - Y_{k-1}, \quad k \geq 1.$$

Then D_k is $\mathcal{F}_k$ -measurable and

$$\mathbb{E}[D_{k+1} \mid \mathcal{F}_k] = \mathbb{E}[Y_{k+1} \mid \mathcal{F}_k] - Y_k = 0 \quad \text{a.s.}$$

Moreover,

$$Y_k - Y_0 = \sum_{i=1}^k D_i.$$

Theorem 1.6.4. Let $\{(D_k, \mathcal{F}_k)\}$ be a martingale difference sequence. Suppose that, for every $k = 1, \dots, n$,

$$\mathbb{E}[e^{\lambda D_k} \mid \mathcal{F}_{k-1}] \leq e^{\lambda^2 \nu_k^2 / 2} \quad \text{for all } |\lambda| < \frac{1}{\alpha_k}.$$

Define

$$\alpha_* := \max_{1 \leq k \leq n} \alpha_k.$$

Then the following statements hold.

(a) $\sum_{k=1}^n D_k$ is sub-exponential with parameters $\left(\sqrt{\sum_{k=1}^n \nu_k^2}, \alpha_*\right)$.

$$(b) \quad \mathbb{P}(|\sum_{k=1}^n D_k| \geq t) \leq \begin{cases} 2 \exp\left(-\frac{t^2}{2 \sum_{k=1}^n \nu_k^2}\right), & 0 \leq t \leq \frac{\sum_{k=1}^n \nu_k^2}{\alpha_*}, \\ 2 \exp\left(-\frac{t}{2\alpha_*}\right), & t > \frac{\sum_{k=1}^n \nu_k^2}{\alpha_*}. \end{cases}$$

Remark 1.6.5. The moment-generating-function assumption in Theorem 1.6.4 is a conditional sub-exponential condition. Part (a) is the martingale-difference analogue of the standard result that a sum of independent sub-exponential random variables is sub-exponential.

Proof. First, the tower property gives

$$\mathbb{E}[D_k] = \mathbb{E}[\mathbb{E}[D_k \mid \mathcal{F}_{k-1}]] = 0,$$

and hence $\mathbb{E}[\sum_{k=1}^n D_k] = 0$.

For $|\lambda| < 1/\alpha_*$, condition successively on the filtration:

$$\begin{aligned} \mathbb{E}\left[e^{\lambda \sum_{k=1}^n D_k}\right] &= \mathbb{E}\left[\mathbb{E}\left[e^{\lambda \sum_{k=1}^n D_k} \mid \mathcal{F}_{n-1}\right]\right] \\ &= \mathbb{E}\left[e^{\lambda \sum_{k=1}^{n-1} D_k} \mathbb{E}\left[e^{\lambda D_n} \mid \mathcal{F}_{n-1}\right]\right] \\ &\leq \mathbb{E}\left[e^{\lambda \sum_{k=1}^{n-1} D_k} e^{\lambda^2 \nu_n^2/2}\right], \quad |\lambda| < \frac{1}{\alpha_n} \\ &= e^{\lambda^2 \nu_n^2/2} \mathbb{E}\left[e^{\lambda \sum_{k=1}^{n-1} D_k}\right], \quad |\lambda| < \frac{1}{\alpha_n} \\ &\leq e^{\lambda^2 \nu_n^2/2} e^{\lambda^2 \nu_{n-1}^2/2} \mathbb{E}\left[e^{\lambda \sum_{k=1}^{n-2} D_k}\right], \quad |\lambda| < \frac{1}{\alpha_n}, |\lambda| < \frac{1}{\alpha_{n-1}} \\ &\leq \dots \leq e^{\lambda^2 \nu_n^2/2} \dots e^{\lambda^2 \nu_1^2/2}, \quad |\lambda| < \frac{1}{\alpha_n}, \dots, |\lambda| < \frac{1}{\alpha_1} \\ &= \exp\left(\frac{\lambda^2}{2} \sum_{k=1}^n \nu_k^2\right), \quad |\lambda| < \frac{1}{\max_{1 \leq k \leq n} \alpha_k}. \end{aligned}$$

This proves part (a). Part (b) follows from the standard two-regime sub-exponential tail bound. $\square$

Remark 1.6.6. If the D_k were independent, the MGF would factor as $\prod_{k=1}^n \mathbb{E}[e^{\lambda D_k}]$. For a martingale difference sequence, conditional expectation replaces this factorization. Given $\mathcal{F}_{n-1}$, the partial sum $\sum_{k=1}^{n-1} D_k$ is $\mathcal{F}_{n-1}$ -measurable and therefore acts as a constant inside the conditional expectation.

1.6.1 Azuma–Hoeffding Inequality

Corollary 1.6.7 (Azuma–Hoeffding random interval). *Let $\{(D_k, \mathcal{F}_k)\}_{k=1}^\infty$ be a martingale difference sequence and suppose $\mathbb{E}[D_k \mid \mathcal{F}_{k-1}] \in [a_k, b_k]$, where $a_k, b_k \in \mathcal{F}_{k-1}$ for all $k = 1, \dots, n$. Assume that $b_k - a_k \leq L_k$ is a constant. Then, for all $t \geq 0$,*

$$\mathbb{P}\left(\left|\sum_{k=1}^n D_k\right| \geq t\right) \leq 2 \exp\left(-\frac{2t^2}{\sum_{k=1}^n L_k^2}\right). \quad (1.61)$$

Proof. Conditioned on $\mathcal{F}_{k-1}$, the random variable D_k still lies in $[a_k, b_k]$ almost surely. By the conditional form of Hoeffding's lemma, D_k is conditionally $(b_k - a_k)/2$ -sub-Gaussian, and therefore $L_k/2$ -sub-Gaussian and is therefore conditionally sub-exponential with parameters

$$\left(\frac{L_k}{2}, 0\right).$$

By Theorem 1.6.4(a),

$$\sum_{k=1}^n D_k \text{ is } \left(\sqrt{\sum_{k=1}^n \frac{(L_k)^2}{4}}, 0\right) \text{ sub-exponential } \longrightarrow \left(\sqrt{\sum_{k=1}^n \frac{(L_k)^2}{4}}\right) \text{ sub-Gaussian.}$$

The stated tail bound follows. $\square$

Corollary 1.6.8 (Azuma–Hoeffding fix interval). *Let $\{(D_k, \mathcal{F}_k)\}$ be a martingale difference sequence. Suppose there are constants $\{(a_k, b_k)\}_{k=1}^\infty$ such that*

$$D_k \in [a_k, b_k] \quad \text{a.s. for every } k.$$

Then, for every $t \geq 0$,

$$\mathbb{P} \left(\left| \sum_{k=1}^n D_k \right| \geq t \right) \leq 2 \exp \left(- \frac{2t^2}{\sum_{k=1}^n (b_k - a_k)^2} \right).$$

Proof. It directly follows from the above corollary. $\square$

Remark 1.6.9. Conditional expectation preserves inequalities: if $X \leq b$, then $\mathbb{E}[X \mid \mathcal{F}] \leq \mathbb{E}[b \mid \mathcal{F}] = b$. Also, every sub-Gaussian random variable is sub-exponential.

1.6.2 Bounded Differences

Definition 1.6.10 (Bounded difference property). We say that a function f satisfies the bounded difference property with parameters $(L_1, \dots, L_n)$ if, for every $k = 1, \dots, n$,

$$|f(x_1, \dots, x_n) - f(x_1, \dots, x_{k-1}, x'_k, x_{k+1}, \dots, x_n)| \leq L_k$$

for all $x_1, \dots, x_n, x'_k$.

Example 1.6.11. For $x, y \in \mathbb{R}^n$, define the Hamming distance by

$$d_H(x, y) := \sum_{i=1}^n \mathbf{1}\{x_i \neq y_i\}.$$

If f is L -Lipschitz with respect to d_H , meaning that

$$|f(x) - f(y)| \leq L d_H(x, y),$$

then changing only the k th coordinate gives

$$|f(x_1, \dots, x_n) - f(x_1, \dots, x_{k-1}, x'_k, x_{k+1}, \dots, x_n)| \leq L \cdot 1 = L.$$

Thus, f has the bounded difference property with $L_k = L$ for every k .

Corollary 1.6.12 (Bounded difference inequality). *Suppose f satisfies the bounded difference property with parameters $(L_1, \dots, L_n)$ and that $X = (X_1, \dots, X_n)$ has independent components. Then, for every $t \geq 0$,*

$$\mathbb{P}(|f(X) - \mathbb{E}[f(X)]| \geq t) \leq 2 \exp \left(- \frac{2t^2}{\sum_{k=1}^n L_k^2} \right).$$

Proof. Let

$$\mathcal{F}_k := \sigma(X_1, \dots, X_k), \quad Y_k := \mathbb{E}[f(X) \mid \mathcal{F}_k].$$

Since $\mathcal{F}_n$ is generated by all the components of X ,

$$Y_n = \mathbb{E}[f(X) \mid \mathcal{F}_n] = f(X),$$

while

$$Y_0 = \mathbb{E}[f(X) \mid \mathcal{F}_0] = \mathbb{E}[f(X)].$$

Therefore,

$$f(X) - \mathbb{E}[f(X)] = Y_n - Y_0 = \sum_{k=1}^n D_k, \quad \text{where } D_k := Y_k - Y_{k-1}.$$

Note

$$\begin{aligned} D_k &= \mathbb{E}[f(X) \mid X_1, \dots, X_k] - \mathbb{E}[f(X) \mid X_1, \dots, X_{k-1}] \\ &= \mathbb{E}_{X_{k+1}, \dots, X_n}[f(X)] - \mathbb{E}_{X_k, X_{k+1}, \dots, X_n}[f(X)] \\ &:= \mathbb{E}_{k+1}[f(X)] - \mathbb{E}_k[f(X)]. \end{aligned}$$

Then $D_k \in \mathcal{F}_k$, and

$$\begin{aligned} \mathbb{E}[D_k \mid \mathcal{F}_{k-1}] &= \mathbb{E}[\mathbb{E}[f(X) \mid \mathcal{F}_k] - \mathbb{E}[f(X) \mid \mathcal{F}_{k-1}] \mid \mathcal{F}_{k-1}] \\ &= \mathbb{E}[f(X) \mid \mathcal{F}_{k-1}] - \mathbb{E}[f(X) \mid \mathcal{F}_{k-1}] \\ &= 0 \quad \text{a.s.} \end{aligned}$$

Thus, $(D_k)_{k=1}^n$ is a martingale difference sequence.

Next,

$$\begin{aligned} D_k &= \mathbb{E}_{k+1}[f(X_1, \dots, X_{k-1}, X_k, X_{k+1}, \dots, X_n)] - \mathbb{E}_k[f(X_1, \dots, X_{k-1}, X_k, X_{k+1}, \dots, X_n)] \\ &\leq \sup_x \mathbb{E}_{k+1}[f(X_1, \dots, X_{k-1}, x, X_{k+1}, \dots, X_n)] - \mathbb{E}_k[f(X_1, \dots, X_{k-1}, X_k, X_{k+1}, \dots, X_n)] \\ &:= B_k, \end{aligned}$$

where $B_k \in \mathcal{F}_{k-1}$. Similarly,

$$\begin{aligned} D_k &\geq \inf_x \mathbb{E}_{k+1}[f(X_1, \dots, X_{k-1}, x, X_{k+1}, \dots, X_n)] - \mathbb{E}_k[f(X_1, \dots, X_{k-1}, X_k, X_{k+1}, \dots, X_n)] \\ &:= A_k, \end{aligned}$$

where $A_k \in \mathcal{F}_{k-1}$. Then, we have

$$\begin{aligned} B_k - A_k &= \sup_x \mathbb{E}_{k+1}[f(X_1, \dots, X_{k-1}, x, X_{k+1}, \dots, X_n)] - \inf_{x'} \mathbb{E}_{k+1}[f(X_1, \dots, X_{k-1}, x', X_{k+1}, \dots, X_n)] \\ &\leq \sup_{x_1, \dots, x_{k-1}, x_k, x', x_{k+1}, \dots, x_n} |f(x_1, \dots, x_k, \dots, x_n) - f(x_1, \dots, x_{k-1}, x'_k, x_{k+1}, \dots, x_n)| \\ &\leq L_k, \end{aligned}$$

using (??). Therefore, by Azuma–Hoeffding with random interval,

$$\mathbb{P}\left(\left|\sum_{k=1}^n D_k\right| \geq t\right) = \mathbb{P}(|f(X) - \mathbb{E}f(X)| \geq t) \leq 2 \exp\left(-\frac{2t^2}{\sum_{k=1}^n L_k^2}\right).$$

□

Example 1.6.13 (U-statistics). Let $g : \mathbb{R}^2 \rightarrow \mathbb{R}$ be a symmetric function of its arguments. Given i.i.d. sequence X_k , $k \geq 1$ of random variables, the quantity

$$U := \frac{1}{\binom{n}{2}} \sum_{1 \leq j < k \leq n} g(X_j, X_k) \quad (1.62)$$

is known as a pairwise U-statistic.

For example, suppose we want to estimate the mean absolute pairwise deviation $\mathbb{E}[|X_1 - X_2|]$. Then U with $g(x, y) = |x - y|$ is one estimate for it.

Since the X_i are identically distributed,

$$\mathbb{E}U = \frac{1}{\binom{n}{2}} \sum_{1 \leq j < k \leq n} \mathbb{E}g(X_j, X_k) = \frac{1}{\binom{n}{2}} \sum_{1 \leq j < k \leq n} \mathbb{E}g(X_1, X_2) = \mathbb{E}g(X_1, X_2).$$

Suppose $\|g\|_\infty \leq b$. Consider $U = f(X_1, \dots, X_n)$. Then

$$\begin{aligned}
& |f(X_1, \dots, X_n) - f(X_1, \dots, X_{k-1}, X'_k, X_{k+1}, \dots, X_n)| \\
&= \left| \frac{1}{\binom{n}{2}} \sum_{1 \leq j < s \leq n} g(X_j, X_s) - \frac{1}{\binom{n}{2}} \sum_{1 \leq j < s \leq n} g(X'_j, X'_s) \right| \quad [\text{where } X'_j = X_j, j = 1, \dots, k-1, k+1, \dots, n], \\
&= \frac{1}{\binom{n}{2}} \left| \sum_{1 \leq j < s \leq n} g(X_j, X_s) - \sum_{1 \leq j < s \leq n} g(X'_j, X'_s) \right| \\
&\leq \frac{1}{\binom{n}{2}} \sum_{1 \leq j < s \leq n} |g(X_j, X_s) - g(X'_j, X'_s)| \\
&\leq \frac{1}{\binom{n}{2}} |2b(n-k+k-1)| \\
&= \frac{4b}{n}.
\end{aligned}$$

So the bounded differences property holds with parameter $L_k = \frac{4b}{n}$ in each coordinate. Therefore, the bounded difference inequality gives

$$\mathbb{P}(|U - \mathbb{E}U| \geq t) \leq 2 \exp\left(-\frac{2t^2}{\sum_{k=1}^n (4b/n)^2}\right) = 2 \exp\left(-\frac{nt^2}{8b^2}\right). \quad (1.63)$$

Example 1.6.14 (Rademacher Complexity). Let $\{\varepsilon_k\}_{k=1}^n$ be an i.i.d. sequence of Rademacher variables. Given a collection of vectors $\mathcal{A} \subset \mathbb{R}^n$, define the random variable

$$Z(\mathcal{A}) = Z := \sup_{a \in \mathcal{A}} \langle a, \varepsilon \rangle = \sup_{a \in \mathcal{A}} \sum_{i=1}^n a_i \varepsilon_i. \quad (1.64)$$

Then $R(\mathcal{A}) := \mathbb{E}Z(\mathcal{A})$ is the Rademacher complexity of $\mathcal{A}$.

Now let us establish $Z(\mathcal{A}) = f(\varepsilon_1, \dots, \varepsilon_n)$ is sub-Gaussian. Given two Rademacher vectors $\varepsilon, \varepsilon' \in \{-1, +1\}^n$, and $\varepsilon'_i = \varepsilon_i$ for $i \neq k$, we have

$$\begin{aligned}
|f(\varepsilon_1, \dots, \varepsilon_n) - f(\varepsilon'_1, \dots, \varepsilon'_n)| &= \left| \sup_{a \in \mathcal{A}} \langle a, \varepsilon \rangle - \sup_{a \in \mathcal{A}} \langle a, \varepsilon' \rangle \right| \\
&\leq \sup_{a \in \mathcal{A}} |\langle a, \varepsilon \rangle - \langle a, \varepsilon' \rangle| \quad (\text{by Lemma 1.6.15}) \\
&= \sup_{a \in \mathcal{A}} |a_k \varepsilon_k - a_k \varepsilon'_k| \\
&= \sup_{a \in \mathcal{A}} |a_k| |\varepsilon_k - \varepsilon'_k| \\
&\leq 2 \sup_{a \in \mathcal{A}} |a_k|.
\end{aligned}$$

Consequently, the bounded difference inequality yields

$$\mathbb{P}(|Z(\mathcal{A}) - R(\mathcal{A})| \geq t) \leq 2 \exp\left(-\frac{2t^2}{\sum_{k=1}^n 4 \sup_{a \in \mathcal{A}} a_k^2}\right) = 2 \exp\left(-\frac{t^2}{2 \sum_{k=1}^n \sup_{a \in \mathcal{A}} a_k^2}\right).$$

Thus, $Z(\mathcal{A})$ is sub-Gaussian with parameter $\sqrt{2 \sum_{k=1}^n \sup_{a \in \mathcal{A}} a_k^2}$.

Lemma 1.6.15. *The absolute difference of the suprema of two functions is bounded by the supremum of their pointwise absolute difference*

$$\left| \sup_{x \in B} f(x) - \sup_{x \in B} g(x) \right| \leq \sup_{x \in B} |f(x) - g(x)|.$$

1.7 Lipschitz Functions of Gaussian Variables

We first show that Lipschitz functions of a standard Gaussian vector are sub-Gaussian.

Definition 1.7.1. A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is L -Lipschitz with respect to the Euclidean norm $\|\cdot\|_2$ if

$$|f(x) - f(y)| \leq L \|x - y\|_2 \quad \text{for all } x, y \in \mathbb{R}^n. \quad (1.65)$$

Theorem 1.7.2 ([1, Theorem 2.26]). Let $X = (X_1, \dots, X_n) \sim N(0, I_n)$, and let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be L -Lipschitz with respect to $\|\cdot\|_2$. Then $f(X) - \mathbb{E}[f(X)]$ is sub-Gaussian with parameter at most L , and hence

$$\mathbb{P}(|f(X) - \mathbb{E}[f(X)]| \geq t) \leq 2e^{-\frac{t^2}{2L^2}} \quad \text{for all } t \geq 0. \quad (1.66)$$

Remark 1.7.3. The sub-Gaussian parameter in Theorem 1.7.2 depends on f only through L and does not depend on the dimension n . Thus an L -Lipschitz function of an n -dimensional standard Gaussian vector obeys the same tail bound as a scalar $N(0, L^2)$ variable, however large n is. The tail bound (1.66) follows from the sub-Gaussian property by applying the Chernoff bound from Lecture 1 to $f(X)$ and to $-f(X)$.

We do not prove Theorem 1.7.2 here; a proof of the weaker statement with sub-Gaussian parameter $\pi L/2$ can be found in [1, Section 2.3]. We now consider three applications.

Example 1.7.4 (Order statistics, [1, Example 2.29]). Given a random vector $(X_1, \dots, X_n)$, its order statistics are obtained by reordering its entries in non-decreasing order,

$$X_{(1)} \leq X_{(2)} \leq \dots \leq X_{(n)}.$$

In particular, $X_{(n)} = \max_{1 \leq i \leq n} X_i$ and $X_{(1)} = \min_{1 \leq i \leq n} X_i$. For $k \in \{1, \dots, n\}$, let $f_k : \mathbb{R}^n \rightarrow \mathbb{R}$ denote the map $x \mapsto x_{(k)}$, so that $X_{(k)} = f_k(X_1, \dots, X_n)$. The Lipschitz constant of f_k is controlled by the following lemma.

Lemma 1.7.5. For all $x, y \in \mathbb{R}^n$ and $k \in \{1, \dots, n\}$,

$$|x_{(k)} - y_{(k)}| \leq \|x - y\|_\infty.$$

The proof of Lemma 1.7.5 is left as Exercise 1. Since $\|x - y\|_\infty = \max_{1 \leq i \leq n} |x_i - y_i| \leq \|x - y\|_2$, the lemma gives

$$|f_k(x) - f_k(y)| \leq \|x - y\|_2 \quad \text{for all } x, y \in \mathbb{R}^n,$$

so each f_k is 1-Lipschitz. Consequently, if $(X_1, \dots, X_n) \sim N(0, I_n)$, Theorem 1.7.2 with $L = 1$ yields, for every $k = 1, \dots, n$,

$$\mathbb{P}(|X_{(k)} - \mathbb{E}[X_{(k)}]| \geq t) \leq 2e^{-\frac{t^2}{2}} \quad \text{for all } t \geq 0. \quad (1.67)$$

In particular, the fluctuations of the maximum $X_{(n)}$ around its mean are of order one, uniformly in n .

Example 1.7.6 (Gaussian complexity, [1, Example 2.30]). Let $A \subseteq \mathbb{R}^n$ be a bounded set and $W \sim N(0, I_n)$, and define

$$Z(A) := \sup_{a \in A} \langle a, W \rangle.$$

The expectation $\mathbb{E}[Z(A)]$ is called the *Gaussian complexity* of A . It is the Gaussian counterpart of the Rademacher complexity from Lecture 6, with the Rademacher vector replaced by W .

To show that $Z(A)$ concentrates, write $Z(A) = f(W)$ with $f(w) := \sup_{a \in A} \langle a, w \rangle$ for $w \in \mathbb{R}^n$, and let $D(A) := \sup_{a \in A} \|a\|_2 < \infty$. For any $w, w' \in \mathbb{R}^n$,

$$\begin{aligned} |f(w) - f(w')| &= \left| \sup_{a \in A} \langle a, w \rangle - \sup_{a \in A} \langle a, w' \rangle \right| \leq \sup_{a \in A} |\langle a, w \rangle - \langle a, w' \rangle| \\ &= \sup_{a \in A} |\langle a, w - w' \rangle| \leq \sup_{a \in A} \|a\|_2 \|w - w'\|_2 = D(A) \|w - w'\|_2. \end{aligned}$$

The second inequality is the Cauchy–Schwarz inequality. Hence f is $D(A)$ -Lipschitz, and Theorem 1.7.2 yields

$$\mathbb{P}(|Z(A) - \mathbb{E}[Z(A)]| \geq t) \leq 2 \exp\left(-\frac{t^2}{2D(A)^2}\right) \quad \text{for all } t \geq 0. \quad (1.68)$$

Example 1.7.7 (Singular values of Gaussian random matrices, [1, Example 2.32]). Let $n > d$ and let $X \in \mathbb{R}^{n \times d}$ be a random matrix with i.i.d. $N(0, 1)$ entries and ordered singular values $\sigma_1(X) \geq \sigma_2(X) \geq \dots \geq \sigma_d(X) \geq 0$. Writing $x_1^\top, \dots, x_n^\top$ for the rows of X , we view the matrix as the vector

$$\text{vec}(X) := (x_1^\top, x_2^\top, \dots, x_n^\top)^\top \in \mathbb{R}^{nd},$$

so that $\text{vec}(X) \sim N(0, I_{nd})$ and each $\sigma_k(X)$ is a function of a standard Gaussian vector in $\mathbb{R}^{nd}$. For $X, Y \in \mathbb{R}^{n \times d}$ we have $\|\text{vec}(X) - \text{vec}(Y)\|_2 = \|X - Y\|_F$, so it suffices to bound $|\sigma_k(X) - \sigma_k(Y)|$ by $\|X - Y\|_F$.

Recall that for $M \in \mathbb{R}^{n \times d}$ the spectral norm and the Frobenius norm are

$$\|M\|_2 := \sigma_1(M) = \sup_{a \in \mathbb{R}^d: \|a\|_2=1} \|Ma\|_2, \quad \|M\|_F := \left(\sum_{i=1}^n \sum_{j=1}^d M_{ij}^2 \right)^{1/2} = (\sigma_1^2(M) + \dots + \sigma_d^2(M))^{1/2},$$

so that $\|M\|_2 \leq \|M\|_F$. By Weyl's inequality (see [1, Exercise 8.3]),

$$\max_{1 \leq k \leq d} |\sigma_k(X) - \sigma_k(Y)| \leq \|X - Y\|_2 \leq \|X - Y\|_F. \quad (1.69)$$

Thus each σ_k is 1-Lipschitz as a function of $\text{vec}(X)$, and Theorem 1.7.2 gives, for each $k = 1, \dots, d$,

$$\mathbb{P}(|\sigma_k(X) - \mathbb{E}[\sigma_k(X)]| \geq t) \leq 2e^{-\frac{t^2}{2}} \quad \text{for all } t \geq 0. \quad (1.70)$$

We have not computed $\mathbb{E}[\sigma_k(X)]$, but (1.70) shows that each singular value stays within $O(1)$ of its expectation with high probability, regardless of n and d .

Chapter 2

Concentration of measure

2.1 Concentration by Entropic Techniques

We now introduce the entropy method for deriving concentration inequalities [1, Section 3.1].

Definition 2.1.1 (ϕ -entropy). Let $X \sim P$, and let ϕ be a convex function on an interval containing the support of X such that X and $\phi(X)$ are integrable. The ϕ -entropy of X is

$$H_\phi(X) := \mathbb{E}[\phi(X)] - \phi(\mathbb{E}[X]). \quad (2.1)$$

Remark 2.1.2. By Jensen's inequality, $H_\phi(X) \geq 0$. Moreover, $H_\phi(X) = 0$ if X is almost surely constant or if ϕ is affine. Although the notation suggests a function of X , $H_\phi(X)$ is a number that depends only on the distribution P of X ; for instance, $H_\phi(e^{\lambda X})$ below is the ϕ -entropy of the distribution of $e^{\lambda X}$.

Example 2.1.3. Two familiar quantities are ϕ -entropies.

- (a) If $\phi(u) = u^2$, then $H_\phi(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2 = \text{Var}(X)$.
- (b) If $\phi(u) = -\log u$ for $u > 0$, then applying ϕ to the positive random variable $e^{\lambda X}$ gives

$$\begin{aligned} H_\phi(e^{\lambda X}) &= \mathbb{E}[-\log e^{\lambda X}] - (-\log \mathbb{E}[e^{\lambda X}]) = -\lambda \mathbb{E}[X] + \log \mathbb{E}[e^{\lambda X}] \\ &= \log e^{-\lambda \mathbb{E}[X]} + \log \mathbb{E}[e^{\lambda X}] = \log \mathbb{E}[e^{\lambda(X - \mathbb{E}[X])}], \end{aligned}$$

which is the cumulant generating function of the centered variable $X - \mathbb{E}[X]$.

From now on we work with the convex function $\phi : [0, \infty) \rightarrow \mathbb{R}$ defined by

$$\phi(u) := u \log u \quad \text{for } u > 0, \quad \phi(0) := 0, \quad (2.2)$$

where the value at 0 makes ϕ continuous, since $u \log u \rightarrow 0$ as $u \rightarrow 0^+$. For a non-negative random variable Z we drop the subscript and write

$$H(Z) := H_\phi(Z) = \mathbb{E}[Z \log Z] - \mathbb{E}[Z] \log \mathbb{E}[Z]. \quad (2.3)$$

Let $\varphi_X(\lambda) := \mathbb{E}[e^{\lambda X}]$ be the moment generating function of X (not to be confused with the convex function ϕ), and assume that it is finite in a neighborhood of λ , so that $\varphi'_X(\lambda) = \mathbb{E}[Xe^{\lambda X}]$. Taking $Z = e^{\lambda X}$ in (2.3) and using $\log e^{\lambda X} = \lambda X$, we obtain

$$H(e^{\lambda X}) = \mathbb{E}[e^{\lambda X} \cdot \lambda X] - \mathbb{E}[e^{\lambda X}] \log \mathbb{E}[e^{\lambda X}] = \lambda \varphi'_X(\lambda) - \varphi_X(\lambda) \log \varphi_X(\lambda). \quad (2.4)$$

Hence $H(e^{\lambda X})$ is determined by the moment generating function of X and its derivative.

Example 2.1.4 (Entropy of a Gaussian variable, [1, Example 3.1]). Let $X \sim N(0, \sigma^2)$. Then $\varphi_X(\lambda) = e^{\lambda^2 \sigma^2 / 2}$ and $\varphi'_X(\lambda) = \lambda \sigma^2 e^{\lambda^2 \sigma^2 / 2}$ for all $\lambda \in \mathbb{R}$, so by (2.4),

$$H(e^{\lambda X}) = \lambda \cdot \lambda \sigma^2 e^{\frac{\lambda^2 \sigma^2}{2}} - e^{\frac{\lambda^2 \sigma^2}{2}} \log e^{\frac{\lambda^2 \sigma^2}{2}} = \frac{\lambda^2 \sigma^2}{2} e^{\frac{\lambda^2 \sigma^2}{2}} = \frac{1}{2} \sigma^2 \lambda^2 \varphi_X(\lambda). \quad (2.5)$$

Since the moment generating function yields tail bounds through the Chernoff bound, (2.4) suggests that an upper bound on $H(e^{\lambda X})$ should also lead to a tail bound. The Herbst argument makes this precise for upper bounds of the Gaussian form (2.5).

Proposition 2.1.5 (Herbst argument, [1, Proposition 3.2]). *Suppose that there is a constant $\sigma > 0$ such that*

$$H(e^{\lambda X}) \leq \frac{1}{2} \sigma^2 \lambda^2 \varphi_X(\lambda) \quad \text{for all } \lambda \in I, \quad (2.6)$$

where $I = \mathbb{R}$ or $I = [0, \infty)$. Then

$$\mathbb{E} \left[e^{\lambda(X - \mathbb{E}[X])} \right] \leq e^{\frac{\lambda^2 \sigma^2}{2}} \quad \text{for all } \lambda \in I. \quad (2.7)$$

Remark 2.1.6. By Example 2.1.4, every $X \sim N(0, \sigma^2)$ satisfies (2.6) with equality for all $\lambda \in \mathbb{R}$. When $I = \mathbb{R}$, the conclusion (2.7) says exactly that X is sub-Gaussian with parameter σ . When $I = [0, \infty)$, it controls the moment generating function only for $\lambda \geq 0$ and so does not by itself give sub-Gaussianity; it does give, through the Chernoff bound from Lecture 1, the one-sided bound

$$\mathbb{P}(X - \mathbb{E}[X] \geq t) \leq e^{-\frac{t^2}{2\sigma^2}} \quad \text{for all } t \geq 0.$$

The proof of Proposition 2.1.5 is deferred; see also [1, Section 3.1.2].

Exercises

1. Prove Lemma 1.7.5: for all $x, y \in \mathbb{R}^n$ and $k \in \{1, \dots, n\}$, $|x_{(k)} - y_{(k)}| \leq \|x - y\|_\infty$.

2.2 new

We begin by completing the discussion of distances and divergences between probability measures. Let $(\mathcal{X}, \rho)$ be a metric space, and let $\mathbb{Q}$ and $\mathbb{P}$ be probability distributions on $\mathcal{X}$.

2.2.1 Optimal Transport Theory

Optimal transport theory studies how to transform one probability distribution $\mathbb{Q}$ into another probability distribution $\mathbb{P}$ while minimizing the cost of transporting probability mass.

A *transportation plan*, or *coupling*, of $\mathbb{Q}$ and $\mathbb{P}$ is a joint probability distribution M of a pair (X, X') such that

$$X \sim \mathbb{Q} \quad \text{and} \quad X' \sim \mathbb{P}.$$

The collection of all couplings of $\mathbb{Q}$ and $\mathbb{P}$ is denoted by $\Pi(\mathbb{Q}, \mathbb{P})$.

Let $\rho(x, x')$ denote the cost of moving one unit of mass from x to x' . For a coupling $M \in \Pi(\mathbb{Q}, \mathbb{P})$, its expected transportation cost is

$$\mathbb{E}_M [\rho(X, X')].$$

Definition 2.2.1 (Wasserstein distance). For $s \geq 1$, the order- s Wasserstein distance is

$$W_{\rho,s}(\mathbb{Q}, \mathbb{P}) := \inf_{M \in \Pi(\mathbb{Q}, \mathbb{P})} (\mathbb{E}_M [\rho(X, X')^s])^{1/s}. \quad (2.8)$$

When $s = 1$, we write $W_{\rho}(\mathbb{Q}, \mathbb{P}) := W_{\rho,1}(\mathbb{Q}, \mathbb{P})$. The case $s = 2$ is also frequently used in practice.

Example 2.2.2 (Transporting bread to coffee stores). Suppose that bread is produced at two store locations, represented by -1 and 1 . The probability measure

$$\mathbb{Q} = \frac{1}{2}\delta_{-1} + \frac{1}{2}\delta_1$$

describes the distribution of the bread supply. Thus, one half of the bread is produced at location -1 , and the other half is produced at location 1 .

The bread must be delivered to coffee stores at the same two locations. The probability measure

$$\mathbb{P} = \frac{1}{3}\delta_{-1} + \frac{2}{3}\delta_1$$

describes the distribution of demand. The coffee store at location -1 requires one third of the total bread, whereas the coffee store at location 1 requires two thirds.

Let $m(x, x')$ denote the amount of bread transported from production location x under $\mathbb{Q}$ to coffee-store location x' under $\mathbb{P}$. The four possible transportation quantities are

$$m(-1, -1), \quad m(-1, 1), \quad m(1, -1), \quad m(1, 1).$$

They have the following interpretations:

- $m(-1, -1)$ is the amount transported from the bread store at -1 to the coffee store at -1 ;
- $m(-1, 1)$ is the amount transported from the bread store at -1 to the coffee store at 1 ;
- $m(1, -1)$ is the amount transported from the bread store at 1 to the coffee store at -1 ;
- $m(1, 1)$ is the amount transported from the bread store at 1 to the coffee store at 1 .

The supply and demand constraints are

$$\begin{aligned} m(-1, -1) + m(-1, 1) &= \frac{1}{2}, & m(1, -1) + m(1, 1) &= \frac{1}{2}, \\ m(-1, -1) + m(1, -1) &= \frac{1}{3}, & m(-1, 1) + m(1, 1) &= \frac{2}{3}. \end{aligned}$$

An optimal transportation plan is

$$\begin{aligned} m(-1, -1) &= \frac{1}{3}, & m(-1, 1) &= \frac{1}{6}, \\ m(1, -1) &= 0, & m(1, 1) &= \frac{1}{2}. \end{aligned}$$

Equivalently, the plan is represented by the table

	Coffee store -1	Coffee store 1	Bread supply
Bread store -1	$m(-1, -1) = \frac{1}{3}$	$m(-1, 1) = \frac{1}{6}$	$\frac{1}{2}$
Bread store 1	$m(1, -1) = 0$	$m(1, 1) = \frac{1}{2}$	$\frac{1}{2}$
Coffee-store demand	$\frac{1}{3}$	$\frac{2}{3}$	1

Suppose that $\rho(x, x') = |x - x'|$. Transporting bread between stores at the same location has zero cost, whereas transporting bread from one location to the other has cost 2. Therefore, the total cost is

$$\begin{aligned} & \rho(-1, -1)m(-1, -1) + \rho(-1, 1)m(-1, 1) \\ & + \rho(1, -1)m(1, -1) + \rho(1, 1)m(1, 1) \\ & = 0 \left(\frac{1}{3} \right) + 2 \left(\frac{1}{6} \right) + 2(0) + 0 \left(\frac{1}{2} \right) = \frac{1}{3}. \end{aligned}$$

The marginal constraints force at least $1/6$ of the bread produced at -1 to be transported to 1 . Hence no coupling can have a smaller cost, and

$$W_\rho(\mathbb{Q}, \mathbb{P}) = \frac{1}{3}.$$

For any 1-Lipschitz function f and any coupling $M \in \Pi(\mathbb{Q}, \mathbb{P})$,

$$\begin{aligned} \int_{\mathcal{X}} f d\mathbb{Q} - \int_{\mathcal{X}} f d\mathbb{P} &= \mathbb{E}_M [f(X) - f(X')] \\ &\leq \mathbb{E}_M [\rho(X, X')]. \end{aligned}$$

Taking the supremum over all 1-Lipschitz functions and then the infimum over all couplings gives the weak-duality inequality.

Theorem 2.2.3 (Kantorovich–Rubinstein duality). *Under the usual regularity and finite-first-moment assumptions,*

$$\begin{aligned} W_\rho(\mathbb{Q}, \mathbb{P}) &= \sup_{\|f\|_{\text{Lip}} \leq 1} \left[\int_{\mathcal{X}} f d\mathbb{Q} - \int_{\mathcal{X}} f d\mathbb{P} \right] \\ &= \inf_{M \in \Pi(\mathbb{Q}, \mathbb{P})} \int_{\mathcal{X} \times \mathcal{X}} \rho(x, x') dM(x, x') \\ &= \inf_{M \in \Pi(\mathbb{Q}, \mathbb{P})} \mathbb{E}_M [\rho(X, X')]. \end{aligned} \tag{2.9}$$

If $\mathbb{Q}$ and $\mathbb{P}$ have densities q and p , and the coupling M has density m , then the transportation cost associated with m is

$$\int_{\mathcal{X} \times \mathcal{X}} \rho(x, x') m(x, x') dx dx'.$$

Thus, the primal formulation (coupling formulation) minimizes transportation cost, whereas the dual formulation (Lipschitz function formulation) compares expectations over 1-Lipschitz functions.

2.2.2 Kullback–Leibler Divergence

Definition 2.2.4 (Kullback–Leibler divergence). The Kullback–Leibler (KL) divergence of $\mathbb{Q}$ from $\mathbb{P}$ is

$$D(\mathbb{Q} \parallel \mathbb{P}) := \begin{cases} \mathbb{E}_{\mathbb{Q}} \left[\log \frac{d\mathbb{Q}}{d\mathbb{P}} \right], & \mathbb{Q} \ll \mathbb{P}, \\ +\infty, & \text{otherwise.} \end{cases} \tag{2.10}$$

If $\mathbb{Q}$ and $\mathbb{P}$ have densities q and p with respect to a common dominating measure ν , then

$$D(\mathbb{Q} \parallel \mathbb{P}) = \int_{\mathcal{X}} q(x) \log \frac{q(x)}{p(x)} \nu(dx). \tag{2.11}$$

Unlike a metric, KL divergence is generally asymmetric and does not satisfy the triangle inequality. It nevertheless controls total variation distance.

Theorem 2.2.5 (Pinsker–Csiszár–Kullback inequality). *With the convention*

$$\|\mathbb{P} - \mathbb{Q}\|_{\text{TV}} := \sup_{A \subseteq \mathcal{X}} |\mathbb{P}(A) - \mathbb{Q}(A)|,$$

we have

$$\|\mathbb{P} - \mathbb{Q}\|_{\text{TV}} \leq \sqrt{\frac{1}{2} D(\mathbb{Q} \| \mathbb{P})}. \quad (2.12)$$

Thus, small KL divergence implies that the two measures assign similar probabilities to every measurable event.

Chapter 3

Uniform Laws of Large Numbers

3.1 Motivation

3.1.1 Uniform Convergence of CDFs

Let X be a real-valued random variable with cumulative distribution function

$$F(t) := \mathbb{P}(X \leq t) = \mathbb{E} [\mathbf{1}_{\{X \leq t\}}], \quad t \in \mathbb{R}.$$

Suppose that

$$X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} \text{Law}(X).$$

Definition 3.1.1 (Empirical CDF). The empirical cumulative distribution function is

$$\hat{F}_n(t) := \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{X_i \leq t\}}, \quad t \in \mathbb{R}. \quad (3.1)$$

For every fixed t , the strong law of large numbers gives

$$\hat{F}_n(t) \longrightarrow F(t) \quad \text{almost surely.}$$

This is pointwise convergence: the exceptional event may initially depend on t .

Many estimation problems can be formulated through a functional γ that maps a CDF F to a real number $\gamma(F)$. The plug-in estimator is $\gamma(\hat{F}_n)$, and the basic consistency question is whether

$$\gamma(\hat{F}_n) \xrightarrow{?} \gamma(F).$$

Example 3.1.2 (Expectation functionals). Let g be an integrable function and define

$$\gamma_g(F) := \int_{\mathbb{R}} g(x) dF(x).$$

Because $\hat{F}_n$ assigns mass $1/n$ to each observation,

$$\gamma_g(\hat{F}_n) = \int_{\mathbb{R}} g(x) d\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n g(X_i).$$

Therefore, if $\mathbb{E} |g(X)| < \infty$, the strong law gives

$$\gamma_g(\hat{F}_n) \longrightarrow \gamma_g(F) \quad \text{almost surely.}$$

Example 3.1.3 (Quantile functionals). For $\alpha \in (0, 1)$, the lower α -quantile functional is

$$Q_\alpha(F) := \inf \{t \in \mathbb{R} : F(t) \geq \alpha\}.$$

The corresponding empirical quantile is

$$\begin{aligned} Q_\alpha(\hat{F}_n) &= \inf \left\{ t \in \mathbb{R} : \hat{F}_n(t) \geq \alpha \right\} \\ &= \inf \left\{ t \in \mathbb{R} : \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{X_i \leq t\}} \geq \alpha \right\}. \end{aligned}$$

It is not immediately clear from pointwise convergence alone whether $Q_\alpha(\hat{F}_n) \rightarrow Q_\alpha(F)$. Uniform convergence and an appropriate condition ensuring that F crosses the level α are needed.

Definition 3.1.4 (Supremum norm and continuity). For two CDFs F and G , their supremum distance is

$$\|F - G\|_\infty := \sup_{t \in \mathbb{R}} |F(t) - G(t)|. \quad (3.2)$$

A functional γ is continuous at F in the supremum norm if, for every $\epsilon > 0$, there exists $\delta > 0$ such that

$$\|G - F\|_\infty \leq \delta \implies |\gamma(G) - \gamma(F)| \leq \epsilon.$$

Recall that if f is a continuous function on $\mathbb{R}^d$, then

$$x_n \rightarrow x \implies f(x_n) \rightarrow f(x).$$

Similarly, if $F_n \rightarrow F$ in the supremum norm and γ is continuous at F , then

$$\gamma(F_n) \rightarrow \gamma(F).$$

Therefore, to establish the consistency of $\gamma(\hat{F}_n)$, it is sufficient to verify that

1. $\|\hat{F}_n - F\|_\infty \rightarrow 0$;
2. γ is continuous at F with respect to the supremum norm.

Theorem 3.1.5 (Glivenko–Cantelli). *For every distribution function F , the empirical CDF $\hat{F}_n$ is a strongly consistent estimator of the population CDF in the supremum norm:*

$$\|\hat{F}_n - F\|_\infty \longrightarrow 0 \quad \text{almost surely.} \quad (3.3)$$

The proof is deferred.

3.1.2 Uniform Laws for General Function Classes

The empirical-CDF result is a special case of a general problem involving empirical averages over a collection of functions. Let

$$X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} \mathbb{P}$$

on a measurable space $\mathcal{X}$.

Definition 3.1.6 (Empirical measure). The empirical probability measure is

$$\mathbb{P}_n := \frac{1}{n} \sum_{i=1}^n \delta_{X_i}. \quad (3.4)$$

For $a \in \mathcal{X}$, the Dirac probability measure δ_a is defined by

$$\delta_a(B) := \begin{cases} 1, & a \in B, \\ 0, & a \notin B, \end{cases} \quad B \subseteq \mathcal{X}. \quad (3.5)$$

Example 3.1.7 (Bernoulli distribution). The Bernoulli distribution with parameter p can be written as

$$\text{Ber}(p) = p\delta_1 + (1 - p)\delta_0.$$

Its probability mass function is

$$\text{pmf}(x) = \begin{cases} p, & x = 1, \\ 1 - p, & x = 0. \end{cases}$$

Thus, if $\mathbb{P} = p\delta_1 + (1 - p)\delta_0$, then

$$\mathbb{P}(\{1\}) = p, \quad \mathbb{P}(\{0\}) = 1 - p.$$

Definition 3.1.8 (Population and empirical averages). For a measurable function f , define

$$\mathbb{P}f := \int_{\mathcal{X}} f(x) d\mathbb{P}(x) = \mathbb{E}_{\mathbb{P}}[f(X)]. \quad (3.6)$$

The corresponding empirical average is

$$\mathbb{P}_n f := \int_{\mathcal{X}} f(x) d\mathbb{P}_n(x) = \frac{1}{n} \sum_{i=1}^n f(X_i). \quad (3.7)$$

For each fixed integrable function f , the strong law of large numbers implies

$$\mathbb{P}_n f \longrightarrow \mathbb{P}f \quad \text{almost surely.}$$

Definition 3.1.9. Let $\mathcal{F}$ be a class of measurable functions. Define

$$\|\mathbb{P}_n - \mathbb{P}\|_{\mathcal{F}} := \sup_{f \in \mathcal{F}} |\mathbb{P}_n f - \mathbb{P}f|. \quad (3.8)$$

The class $\mathcal{F}$ is called a $\mathbb{P}$ -Glivenko–Cantelli class if

$$\|\mathbb{P}_n - \mathbb{P}\|_{\mathcal{F}} \longrightarrow 0 \quad \text{almost surely.} \quad (3.9)$$

Example 3.1.10 (Empirical CDFs and indicator functions). Define the threshold class

$$\mathcal{F}_{\text{cdf}} := \{f_t(\cdot) : t \in \mathbb{R}\}, \quad f_t(x) := \mathbf{1}_{\{x \leq t\}}.$$

For every $t \in \mathbb{R}$,

$$\mathbb{P}f_t = \int_{\mathcal{X}} \mathbf{1}_{\{x \leq t\}} d\mathbb{P}(x) = \mathbb{P}(X \leq t) = F(t),$$

and

$$\mathbb{P}_n f_t = \frac{1}{n} \sum_{i=1}^n f_t(X_i) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{X_i \leq t\}} = \hat{F}_n(t).$$

Therefore,

$$\|\mathbb{P}_n - \mathbb{P}\|_{\mathcal{F}_{\text{cdf}}} = \sup_{t \in \mathbb{R}} |\hat{F}_n(t) - F(t)| = \|\hat{F}_n - F\|_{\infty}.$$

The Glivenko–Cantelli theorem says precisely that the threshold class $\mathcal{F}_{\text{cdf}}$ is a universal Glivenko–Cantelli class.

This reformulation is the bridge from empirical CDFs to empirical-process theory. For a single function, the law of large numbers is enough. For an entire class, the size and structure of the class determine whether the same sample controls all empirical averages simultaneously.

Bibliography

- [1] Martin J. Wainwright, *High-Dimensional Statistics: A Non-Asymptotic Viewpoint*, Cambridge University Press, 2019.
- [2] R. Vershynin, *High-Dimensional Probability: An Introduction with Applications in Data Science*, Cambridge University Press, 2018.